AFRICA GOVERNANCE AND CIVIC INNOVATION HUB (AGCIH) POLICY COMMENTARY

Kenya Artificial Intelligence Bill, 2026:

Governance Architecture, Administrative Capacity, and the Rule-of-Law Threshold

By Danai Hazel Kudya
Founder and Executive Director
Africa Governance and Civic Innovation Hub (AGCIH)
March 2026

AGCIH does not advocate, regulate, or deploy technology. AGCIH operationalises governance.

Document Type	Policy Commentary
Subject Bill	The Artificial Intelligence Bill, 2026 (Senate Bills No. 4)
Jurisdiction	Republic of Kenya
AGCIH Series	Governance Architecture and AI Regulation

Executive Summary

The Kenya Artificial Intelligence Bill, 2026, published on 19th February 2026, represents one of the most substantive AI governance proposals currently before an African legislature. It establishes an Office of the Artificial Intelligence Commissioner, adopts a risk-based classification framework, introduces regulatory sandboxes, mandates transparency obligations, and creates criminal liability for specified violations. As a foundational legislative framework, it is credible.

This commentary does not dispute the Bill's ambition. It tests the Bill's implementability: whether the institutional architecture it creates can, in practice, sustain the governance functions it mandates on rule of law terms.

AGCIH's analysis identifies four structural gaps that, if not addressed before the Bill's commencement, will produce a governance framework that exists on paper but operates unevenly or harmfully in practice.

GAP 1 — Administrative Hosting Capacity Deficit

The Office of the Artificial Intelligence Commissioner is granted an extraordinarily broad operational mandate encompassing risk assessments, conformity audits, sandbox management, nationwide literacy programmes, and enforcement without any phased operationalisation provision or minimum governance readiness threshold that must be met before enforcement powers are activated.

GAP 2 — Procurement Entry Failure

Section 34, the Bill's entire provision on public sector AI use, contains a single clause. It requires public entities to comply with the Act but creates no procurement governance framework, no pre-deployment assessment requirement, no internal responsibility architecture, and no documentation standard. AI enters public administration primarily through procurement; this is not a minor drafting omission. It leaves the state's principal point of AI exposure substantially ungoverned.

GAP 3 — Criminal Liability Before Standards

The offences and penalties under Part VI are in effect upon commencement, but the standards against which compliance is measured, risk classification criteria, assessment procedures, and mitigation requirements are deferred to subordinate regulations yet to be written. Criminal liability precedes the standards it enforces. This is a rule of law design problem, not a technical drafting issue.

GAP 4 — Relocation of Judgment Without Anchoring

Section 32 mandates human oversight in critical AI decisions but defers the definition of 'critical decisions' to Cabinet Secretary regulations. Until those regulations exist, the human oversight protection is formally present but operationally hollow. This is what AGCIH identifies as Relocation of Judgment risk: accountability appears assigned but cannot, in practice, be located or enforced.

This commentary advances six specific recommendations for legislative amendment. Each recommendation is actionable within the Bill's existing structure and does not require the Bill to be redrafted. The recommendations are addressed to the Senate, to the Office of the Attorney-General, and to the technical working processes surrounding the Bill's committee stage.

AGCIH's analysis draws on its doctrinal framework of Continuous Administration, Administrative Hosting Capacity, Relocation of Judgment, Procurement Entry Doctrine, Governance Readiness, and Institutional Friction developed through AGCIH's sustained work on AI governance and public institutional design in the African context.

1. Introduction: The Governance Moment Kenya Must Get Right

Kenya is moving faster on AI governance than most African states. Its National AI Strategy (2025–2030) was launched in March 2025. The AI Bill was introduced in February 2026, reportedly accelerated by a High Court order following an urgent petition. A Commissioner's appointment process is anticipated to follow. The legislative momentum is real and, in terms of regional standard-setting, consequential.

That momentum creates its own risk. The governance conversation has been dominated by two framings: the ethics and equity framing, which asks whether the Bill protects the vulnerable and promotes fair access; and the capacity observation framing, which notes that effective enforcement requires skills and resources that are difficult to assemble. Both observations are accurate. Neither is sufficient.

The question this commentary asks is different: is the Bill designed to be implementable on rule of law terms? Not merely capable of good outcomes in ideal conditions, but structured so that the institutions it creates can actually perform the governance functions it mandates accountably, reviewably, and without producing new forms of administrative harm in the process of regulating AI-related harm.

This is the governance architecture question. It is before both ethics and capacity. A Bill can be ethically well intentioned and technically expert in its drafters, and still produce an institutional design that, upon commencement, creates obligations no institution can discharge, standards no one can apply, and liability regimes that operate before the rules they enforce have been written.

The Kenya AI Bill, 2026, has elements of precisely that problem. This commentary names them with specificity and proposes the amendments required to resolve them.

2. What the Bill Does Well

Before identifying its structural gaps, it is important to characterise the Bill accurately. Several of its features are genuinely well designed and, in the African legislative context, represent forward movement.

2.1 The Risk-Classification Architecture

Part V's four-tier risk classification framework — unacceptable risk, high risk, limited risk, and minimal risk is the right structural approach. It avoids both overregulation of low-risk applications and underregulation of systems with serious societal implications. The inclusion of public administration as a high-risk sector in Section 25(2)(b) is particularly significant and correct. It signals that government AI is not exempt from rigorous oversight by virtue of its origin.

2.2 The Transparency and Disclosure Obligations

Section 28's transparency requirements disclosure of the nature and limitations of a system, the extent of automated decision making, and bias mitigation measures are substantive obligations, not aspirational language. The requirement that high-risk providers submit annual compliance reports (Section 28(3)) creates a routine accountability mechanism that is administrable in practice.

2.3 The Workforce Impact Obligations

Section 33's requirement that deployers of employment impacting AI conduct workforce impact assessments and implement mitigation measures, including reskilling programmes in collaboration with the government, is an uncommon and commendable feature in AI legislation globally, and particularly apt for the Kenyan labour market context.

2.4 The Advisory Committee Composition

Section 17's Advisory Committee is well constructed. Its membership includes civil society representation, Data Protection Commissioner representation, science and technology institutional presence, and gender balance requirements. The consultative nomination processes for private sector and civil society members, through registered organisations rather than ministerial appointment, provide a structural buffer against capture.

2.5 The Three-Year Review Requirement

Section 37's mandatory review every three years, with findings submitted to Parliament, builds in a self-correction mechanism that recognises AI regulation cannot be static. This is governance literate drafting. These features are real strengths. The commentary that follows does not diminish them. It examines what must be corrected for these strengths to function as intended.

3. The Four Structural Gaps

AGCIH's analysis identifies four points in the Bill's architecture where the governance design fails to match the governance ambition. Each gap is identified by its location in the Bill, its precise nature, and its rule of law implications.

3.1 Gap One: The Commissioner's Office: Administrative Hosting Capacity Deficit

The Office of the Artificial Intelligence Commissioner is granted, under Section 10, one of the broadest regulatory mandates in Kenyan statutory law. Its functions include: overseeing implementation and enforcement of the Act; conducting risk assessments; performing conformity audits and post-market surveillance; assessing high-risk systems for compliance; developing policies, guidelines, codes of practice and standards; managing regulatory sandboxes; receiving and investigating complaints; collaborating with national, regional and international bodies; advising national and county governments; promoting public awareness; developing capacity-building programmes at national and county levels; conducting research; maintaining a public register of high-risk systems; and promoting equitable access to AI infrastructure.

This mandate is not advisory. It is operational. It requires, at minimum: specialised technical staff capable of conducting algorithmic audits and conformity assessments; legal and administrative staff capable of processing complaints, issuing enforcement notices, and managing appeals; policy staff capable of producing codes, guidelines, and standards across multiple sectors; outreach and education capacity capable of operating at county level; research capacity; and information management systems capable of maintaining a public register and processing annual compliance reports from high-risk system providers across the economy.

The Bill provides for the appointment of deputy commissioners, assistant commissioners, and other staff under Section 12. It creates financial provisions under Part IV. These are necessary but not sufficient. What the Bill does not provide is a phased operationalisation framework; a requirement that the Office demonstrates by reference to defined minimum governance criteria, that it is institutionally ready to exercise particular functions before those functions are activated.

> **The AGCIH Doctrine of Administrative Hosting Capacity**
>
> *An institution must be capable of hosting a governance function before it can legitimately exercise it. Hosting means more than formal establishment. It means the institution can define boundaries, understand scope, supervise use, investigate failures, retain records, and explain its actions to affected parties and to the public. A regulator that cannot host its own mandate becomes a governance structure that generates legal exposure without generating legal protection.*

The consequence of this gap is predictable. The Office will commence operations with its full mandate legally active and its full enforcement powers legally available. It will face immediate demands, complaints, registration requirements, audit requests, and county government enquiries before it has staffed, systemised, or operationally grounded the functions those demands require. Early enforcement action taken without adequate institutional capacity will be legally fragile. Early inaction in the face of legitimate complaints will erode public credibility. Neither outcome serves the Bill's objectives.

3.2 Gap Two: Section 34 — The Procurement Entry Failure

Section 34 is the Bill's entire provision governing AI use by public institutions. It reads, in full:

> *"34. A public entity, including a county government, that uses an artificial intelligence system shall ensure compliance with this Act."*

That is the entirety of Section 34. One sentence.

The section does not specify who within a public entity holds responsibility for compliance. It does not require a pre-deployment governance assessment. It does not address AI procurement; the moment at which public institutions are most exposed to capability-exceeding or governance-undermining systems. It does not require documentation of AI use within institutional workflows. It does not mandate internal supervision protocols, audit trails, or staff training as conditions of lawful deployment. It does not differentiate between a national ministry deploying a high-risk AI system in a criminal justice context and a county government using a scheduling tool. This is not a minor drafting omission. It leaves the state's principal point of AI exposure substantially ungoverned.

The Bill's risk-based framework under Part V, with its detailed obligations for high-risk deployers in the private sector, simply does not extend in any meaningful way to the public sector. The private sector deployer of a high-risk system faces obligations under Sections 26, 27, 28, and 29. The public sector deployer faces Section 34: comply with the Act.

> **The AGCIH Procurement Entry Doctrine**
>
> *AI enters public administration primarily through procurement. The governance risk is highest at the moment of institutional entry when a public body decides to acquire, integrate, or authorise a system within its administrative workflows. Post-deployment enforcement cannot substitute for governance at the entry point. By the time a regulator investigates a public institution's AI-related harm, the harm is already operational and often structurally embedded in administrative practice.*

This is the most consequential gap in the Bill. Public sector AI use in healthcare administration, social services, tax assessment, land management, court administration, immigration, public employment, and procurement itself carries some of the highest stakes for citizens. The Bill's framework for governing that use consists of one sentence that adds nothing to what the Act's general provisions already require.

A governance framework that applies less rigour to the public sector than to private deployers inverts the accountability structure that the rule of law governance requires. Public institutions exercise public authority. Their use of AI is subject to constitutional obligations of fairness, accountability, and transparency that private actors do not carry in the same form. The Bill's public sector governance architecture must reflect that distinction.

3.3 Gap Three: The Regulations Architecture: Criminal Liability Before Standards

Section 35 (Offences and Penalties) is active upon the Bill's commencement. Under Section 35(1), a person commits an offence if they, among other things, deploy a high-risk AI system without conducting the required risk assessment or implementing the mitigation measures prescribed under Section 26; or fail to comply with disclosure or transparency obligations under Section 28.

The penalties for these offences include fines of up to five million shillings and imprisonment of up to two years.

Section 36(2) reads, “The regulations to be made by the Cabinet Secretary shall provide for, among other things: the detailed criteria, examples, and procedures for classifying AI systems according to risk levels; the forms, processes, and timelines for conducting risk assessments and conformity audits for high-risk systems; and the content, format, and review processes for transparency disclosures and compliance reports.”

The offences under Section 35 are based on non-compliance with requirements whose operational detail asks what risk classification means in practice, what form a risk assessment must take, what a transparency disclosure must contain, which are to be determined by regulations that have not yet been written and are not required to be written before commencement.

This creates a direct rule of law problem. The principle of legal certainty, recognised in the Constitution of Kenya under Article 47 (fair administrative action) and in the broader framework of the rule of law under Article 10, requires that persons subject to law must be able to know in advance what the law requires of them. A criminal provision that activates before the standards of compliance it enforces have been defined does not satisfy that requirement.

The Governance Readiness Principle

Governance readiness means that the institutional and regulatory infrastructure necessary to exercise a power, enforce an obligation, or apply a standard must be in place before that power is exercised, that obligation is enforced, or that standard is applied. A Bill that creates criminal liability for non-compliance with standards yet to be defined is institutionally incomplete at commencement.

This is not a minor procedural concern. Deployers, including public institutions, will be subject to criminal exposure for non-compliance with requirements whose operational meaning does not yet exist in law. Legal advisers will be unable to provide clients with reliable guidance on the conduct required. The Commissioner, if appointed before regulations are finalised, will be in the position of enforcing obligations that it has not yet defined.

3.4 Gap Four: Section 32: Relocation of Judgment Without Anchoring

Section 32 is the Bill's human-centric AI provision. It requires that a person who designs or deploys an AI system must design or deploy it in a manner that enhances rather than replaces human capabilities; incorporate features that support human involvement; and provide for human oversight in critical decisions made by the system. Section 32(2) specifies that oversight in critical decisions must include review mechanisms that allow a qualified person to intervene or override the system's outputs where decisions may affect human rights, safety, or societal well-being.

This is the right principle. Human oversight of AI in high-stakes decisions is a core governance protection.

The problem is Section 32(3). The Cabinet Secretary shall make regulations prescribing, among other things, the critical decisions requiring human oversight. The definition of what constitutes a “critical decision” and therefore the scope of the human oversight obligation is deferred to regulations. Until those regulations are made, a deployer can argue in good faith that their system's decisions do not fall within the undefined category of ‘critical.' The human oversight protection exists as text but not as an enforceable standard. No deployer can be held to a requirement whose scope has not been defined. No affected person can rely on an oversight mechanism whose trigger conditions are unspecified.

> **The AGCIH Doctrine of Relocation of Judgment**
>
> *Relocation of Judgment occurs when formal authority is nominally retained by a human institution or officer, while operational reality reorganises decision-making around system outputs. The risk is not only formal delegation, but also the condition in which accountability appears assigned but cannot, in practice, be located, exercised, or enforced. A provision that guarantees human oversight without defining when that oversight is required is a provision that permits the relocation of judgment while appearing to prevent it.*

The definition of 'critical decisions' is a governance question, not a technical one. It involves choices about which categories of state or institutional action are too consequential to be shaped by automated systems without mandatory human review. That choice belongs in primary legislation, not in subordinate instruments that may be delayed, contested, or quietly circumscribed.

4. Summary: The Four Gaps at a Glance

Gap	Bill Location	Governance Problem	Rule of Law Risk
Administrative Hosting Capacity	Part II — Office establishment; no phased operationalisation	The full mandate is active before the Office can institutionally perform it	Enforcement without capacity; public credibility erosion
Procurement Entry	Section 34 — single sentence	No public sector procurement governance, no pre-deployment assessment	Government AI is ungoverned at the point of entry; less rigour than in the private sector
Regulations Architecture	Sections 35 and 36	Criminal liability is active before compliance standards are defined	Violation of legal certainty; unenforceable on commencement
Relocation of Judgment	Section 32(3)	'Critical decisions' requiring human oversight are deferred to unwritten regulations.	Human oversight protection formal but unenforceable

5. Recommendations

The following six recommendations address each structural gap identified in this commentary. Each recommendation is specific, actionable within the Bill's existing structure, and does not require a full redraft. They are addressed to the Senate, to the Office of the Attorney-General and Department of Justice, and to any technical working process engaged in the Bill's committee stage.

Recommendation 1: Insert a Phased Operationalisation Provision for the Commissioner's Office

The Bill should be amended to include a phased operationalisation provision under Part II, requiring the Commissioner to publish a Governance Readiness Plan within six months of appointment. The Plan must demonstrate, for each category of function under Section 10, the staffing, technical systems, administrative protocols, and minimum institutional capacity that the Office has established or has a funded timeline to establish before exercising that function.

Specifically, the Bill should differentiate between functions the Office may exercise from the date of the Commissioner's appointment (receipt of complaints, publication of the high-risk register, advisory functions) and functions that require demonstrated operational readiness before activation (enforcement action, conformity audits, imposition of administrative fines). This phasing is not a limitation on the Office's mandate. It is the condition that makes the mandate credible.

Recommendation 2: Replace Section 34 with a Public Sector AI Governance Framework

Section 34 must be replaced with a substantive provision governing AI use in the public sector. At a minimum, the replacement provision should require:

- that each public entity designates a named Accountable Officer for AI governance before deploying any AI system in its administrative operations;
- that any public entity intending to procure or deploy an AI system classified or likely to be classified as high-risk must notify the Commissioner and complete a pre-deployment governance assessment, the minimum content of which is defined in primary legislation;
- that procurement contracts for AI systems by public entities include specified governance provisions covering transparency, suspension rights, audit access, and vendor obligations in the event of system failure or harm to citizens;
- that public entities maintain documentation of AI system use sufficient to support internal review, external audit, and the exercise of rights by affected persons; and
- that the Commissioner publishes and periodically updates a Public Sector AI Governance Code setting out the standards against which public entity compliance is assessed.

The replacement provision should make explicit that public entities are subject to the full obligations applicable to high-risk deployers under Part V where they deploy systems falling within the high-risk classification.

Recommendation 3: Establish a Regulations-First Commencement Condition for Enforcement Provisions

The offences and penalties under Part VI, Sections 35 and 36, should not come into force until the Cabinet Secretary has gazetted regulations prescribing, at minimum: the detailed criteria for risk classification under Section 25; the forms, processes, and timelines for risk assessments and human rights impact assessments under Section 26; and the content and format requirements for transparency disclosures under Section 28.

This may be achieved by inserting a commencement provision specifying that Part VI comes into force on a date to be notified by the Cabinet Secretary in the Gazette, and that no such notice may be issued before the regulations required by Sections 25, 26, and 28 are in force. Alternatively, a savings clause should be inserted providing that no criminal proceedings may be commenced under Section 35 in respect of conduct for which the applicable compliance standard has not yet been defined by regulations.

The principle at stake is straightforward: criminal liability requires prior notice of the conduct it criminalises. That notice cannot be given by a Bill alone when the Bill expressly defers the operative standard to regulations yet to be written.

Recommendation 4: Define 'Critical Decisions' in Primary Legislation

Section 32's human oversight obligation should not be contingent on regulations for its operational definition. The Bill should define, in primary legislation, a non-exhaustive list of decision categories that constitute 'critical decisions' for Section 32, including decisions affecting liberty, welfare entitlements, employment, housing, health treatment, immigration status, criminal justice outcomes, and access to public services.

The Cabinet Secretary may, by regulation, extend this list. But the baseline must be in the Act. The protection of human oversight in decisions affecting fundamental rights is not a matter appropriate for subordinate instruments subject to ministerial discretion and potentially delayed, revised, or withdrawn without primary legislative action.

Recommendation 5: Introduce a Regulatory Readiness Assessment for the Commissioner's Appointment

The qualification requirements for the Artificial Intelligence Commissioner under Section 6 are robust in their technical and governance dimensions. The Bill should additionally require that the Commissioner, within ninety days of taking office, publish a public institutional readiness assessment identifying the governance functions the Office is prepared to exercise, those requiring further resourcing, and a timeline for full operationalisation. This assessment should be submitted to the relevant Parliamentary Committee and made publicly available.

This is not a burden on the Commissioner. It is transparency about the institution's actual starting condition, and it creates the public accountability that the Office will need to function with legitimacy from its earliest days.

Recommendation 6: Clarify the Relationship Between the Commissioner and Existing Sectoral Regulators

The Bill's mandate extends across healthcare, education, agriculture, finance, security, employment, and public administration, each of which has an existing regulatory body. The Bill does not specify how the Commissioner's authority relates to, is coordinated with, or is bounded by the authority of the

Communications Authority, the Data Protection Commissioner, the Capital Markets Authority, the Kenya Revenue Authority, or sector-specific regulators.

This is not merely an administrative coordination issue. It is a question of legal certainty. A high-risk AI system deployed in a licensed financial services context is subject to both the Commissioner's requirements under this Bill and the regulatory framework of the Capital Markets Authority or Central Bank. Which standard prevails in the event of conflict? Who conducts the conformity audit? Where does a complaint lie?

The Bill should include a provision establishing a mandatory coordination protocol between the Commissioner and named sectoral regulators, specifying lead authority, concurrent obligations, and conflict resolution mechanisms. Clarity on regulatory jurisdiction is not a luxury in complex institutional environments. It is the condition that makes enforcement coherent.

6. The AGCIH Governance Framework: Why Architecture Comes First

The analysis in this commentary is grounded in AGCIH's doctrinal framework for AI governance in public administration. Each doctrine is named and briefly characterised below. Together, they provide a coherent analytical vocabulary for the governance architecture problems that AI legislation in African contexts must address.

AGCIH Doctrine	Application to the Kenya AI Bill
Continuous Administration	Public institutions do not govern episodically. They administer continuously. AI systems enter this continuous environment and alter its character over time, not at a single point. Governance cannot be treated as a one-time compliance check.
Administrative Hosting Capacity	A public institution or regulator must be able to host an AI system, meaning it can define boundaries, supervise use, investigate failures, and explain its actions before that system is operationally authorised. The Commissioner's Office Gap is a hosting capacity deficit.
Procurement Entry Doctrine	The governance risk for public institutions is highest at the procurement decision. Section 34's failure to govern public sector AI procurement leaves the entry point, the most consequential governance moment, unaddressed.
Relocation of Judgment	When accountability appears assigned but cannot in practice be located, exercised, or enforced, judgment has been relocated without authorisation. Section 32's undefined 'critical decisions' creates this condition.
Governance Readiness	Readiness is not ambition. It is an institutional condition, the presence of defined responsibility, supervision protocols, documentation standards, and contestation pathways before a governance function is exercised. The Regulations Architecture Gap is a readiness failure.
Institutional Friction	AI does not enter neutral environments. It meets legal procedures, professional duties, constitutional obligations, and public expectations. Where governance architecture does not prepare institutions for these friction points, frictions accumulate silently as over-reliance, poor documentation, or decisions that appear attributable but are system-shaped.

These doctrines are not abstract principles. They are diagnostic tools. Applied to the Kenya AI Bill, they identify specific provisions that must be strengthened, specific gaps that must be closed, and specific design choices that, if left unaddressed, will produce governance failures that no amount of enforcement capacity can later repair.

7. Conclusion

The Kenya Artificial Intelligence Bill, 2026, is a serious legislative effort. It names the right problems, reaches toward the right principles, and creates a structural framework capable of genuine governance if that framework is completed.

The four gaps identified in this commentary are not peripheral. They go to the centre of whether the Bill, as currently drafted, can produce the governance outcomes it seeks. A regulator without a hosting capacity pathway will generate either overreach or paralysis. A public sector governance provision of one sentence will leave government AI ungoverned at the point of greatest risk. A criminal liability regime that activates before compliance standards exist will generate legal uncertainty without accountability. A human oversight protection that cannot be triggered because its scope is undefined will relocate judgment while appearing to anchor it.

The recommendations in this commentary are directed at resolving these four gaps. They do not ask Parliament to redraft the Bill. They ask Parliament to complete it to ensure that the ambition the Bill expresses is matched by the institutional architecture required to realise it.

Kenya has an opportunity to enact AI governance legislation that the continent can learn from. That opportunity is best taken by getting the architecture right before commencement, not by discovering its gaps after the Commissioner is in office, the enforcement machinery is operating, and the public sector is using systems whose governance framework consists of a single sentence.

The rule of law is not preserved by ambition. It is preserved by design. And in AI governance, the most important design question is this: can the institutions this legislation creates remain fully responsible for the authority they are given to exercise in public?

AGCIH submits that, with the amendments proposed in this commentary, the answer can be yes.

About AGCIH

The Africa Governance and Civic Innovation Hub (AGCIH) is an independent governance institution based in Harare, Zimbabwe. AGCIH works at the intersection of administrative law, digital governance, and institutional design, with a focus on AI and digital systems in African public administration.

AGCIH is positioned as a governance institution and institutional architect, not a civil society organisation, a technology company, or a policy consultancy. AGCIH's work is grounded in the conviction that Africa must develop its own governance frameworks for AI, built from its own institutional realities, rather than transplanting frameworks designed for different administrative contexts.

AGCIH's doctrinal framework, comprising the doctrines of Continuous Administration, Administrative Hosting Capacity, Relocation of Judgment, Administrative Risk, Institutional Absorption, Procurement Entry Doctrine, Functional Drift, Institutional Friction, Pro-Governability, and Governance Readiness, has been developed through sustained analytical work on AI governance in the African public sector.

AGCIH does not advocate, regulate, or deploy technology. AGCIH operationalises governance.

Contact and Correspondence
Africa Governance and Civic Innovation Hub (AGCIH)
Harare, Zimbabwe
admin@agcih.africa
www.agcih.africa
Founder and Executive Director: Danai Hazel Kudya